

# Power through Simulation

Greg Snow

February 8, 2018

## Motivation

### Power

Two common questions in study design:

- How big of a sample do I need?
- What are my chances of a successful study?

### Other Questions

Questions that should be asked more:

- What does *No Effect* look like
- What does *low power* look like
- Type I, Type II, Type S, and Type M errors
- Robustness to assumptions and their effect on power

## Power Training

- Simple cases
  - One sample z-test with known variance
  - One sample test of proportions
- Pre Built tools

## Pre Built Tools

- Simple problems only
- Complex/confusing Effect Size
- Hidden Assumptions
- Not Flexible Enough

# Power


## What is needed for Power

- Sample Size
  - Balanced/Unbalanced
  - Numbers of Nested Units
- Measure(s) of Variability
- Expected or Minimally Interesting Difference (Effect Size)
- Other assumptions, covariates, etc.

## What is needed for Sample Size

- Desired Power(s)
- Measure(s) of Variability
- Expected or Minimally Interesting Difference (Effect Size)
- Other assumptions, covariates, etc.

## Confidence Interval Approach



## Examples

### Examples

The following are problems where existing tools do not apply or were not satisfactory.

### Non-normal Data

The Central Limit Theorem lets us use the t-test and regression when the population is not normal, but the sample size is large enough. But what is large enough and what effect does the non-normality have on the power?

### Fishers Exact Test

Small sample contingency table where  $\chi^2$  analysis is unlikely to be accurate. Plan to use Fishers Exact test.

### Binomial CI

Study was to show that a rare event had a rate under 4% (previous data suggested about 2%). Planned to find 95% CI using Binomial with a uniform prior and show that entire CI is below 4%

### Regression

Full-Reduced Regression model. Do  $x_1$ - $x_3$  have a significant effect after accounting for  $x_4$ - $x_8$ ? With a few different assumed covariance structures.

### Survival Analysis

Cox Proportional Hazards model with censored and truncated data. Compare different assumptions on amount of censoring.

### Mixed Effects Model

Testing main effect and interaction on students who are nested in schools (Knowledge and Fitness are both outcomes).

### Non-Inferiority

Non-inferiority studies declare success as long as the confidence interval on the parameter of interest is greater than an equivalence value.

### Multivariate Response

When you have multiple outcome variables then tools like Wilks Lambda can be appropriate. But existing tools are only for univariate outcomes.

## Power Through Simulation

### 4 Steps

1. Decide what your data is going to look like
2. Decide how you will analyze the data
3. Simulate data from 1, then analyze it using 2
4. Repeat 3 a bunch of times

The power is the proportion of times that the null is rejected.

(You may iterate a few times between 1 and 2)

### Data

What will your data look like?

- What things are you likely to change?
  - Sample Size
  - Effect Size (difference between means)
  - Regression Coefficients
- Normal or other distribution(s)?
- What will be constant? (or does not matter)
  - Mean of group 1
  - Intercept
  - Variance/Standard Deviation
- Shape/structure of data

### Analysis

- How will you analyze the data?
- What test will you use?
- What is your cut-off for significance?

### Simulation

Write a function that will generate the data and analyze it.

Things that you will change should be function arguments.

### Run the simulation a bunch of times

- Parallel processing if available.
- Run with the Null true to verify Type I error rate.
- Start with 100 runs until you find the parameter values of interest.
- Run 10,000 to 1,000,000 runs for the conditions of interest.

## Simulation Examples

### Starting with what we know

```
> power.t.test(n=300, delta=0.8, sd=3)
```

Two-sample t test power calculation

```
      n = 300
  delta = 0.8
      sd = 3
sig.level = 0.05
  power = 0.9033319
alternative = two.sided
```

NOTE: n is number in *each* group

### Simulating in R

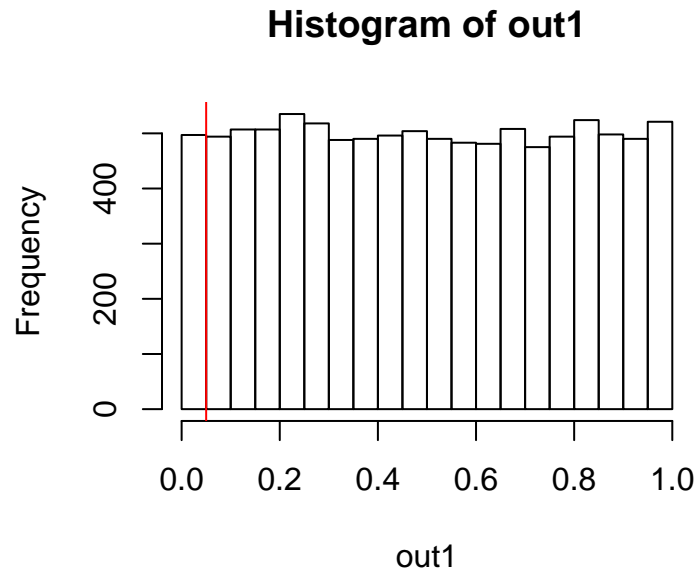
```
> simfun <- function(n=100, diff=0, sd=1) {
+   x1 <- rnorm(n, 0, sd)
+   x2 <- rnorm(n, diff, sd)
+   t.test(x1, x2)$p.value
+ }
```

### Simulating under the Null

```
> out1 <- replicate(10000, simfun(n=300, diff=0.0, sd=3))
> mean(out1 <= 0.05)
```

```
[1] 0.0497
```

```
> hist(out1)
> abline(v=0.05, col='red')
```




#### Simulating for Power

```
> out2 <- replicate(10000, simfun(n=300, diff=0.8, sd=3))  
> mean(out2 <= 0.05)
```

```
[1] 0.8987
```

```
> hist(out2)  
> abline(v=0.05, col='red')
```

## Histogram of out2



### Confidence Interval on the Power

```
> binom.test(sum(out2<=0.05), length(out2))
```

Exact binomial test

```
data: sum(out2 <= 0.05) and length(out2)
number of successes = 8987, number of trials
= 10000, p-value < 2.2e-16
alternative hypothesis: true probability of success is not equal to 0.5
95 percent confidence interval:
 0.8926197 0.9045470
sample estimates:
probability of success
      0.8987
```

### Another Approach

```
> diff <- 0.8
> n <- 300
> sd <- 3
> nsim <- 10000
```

```

> x1 <- matrix(rnorm(n*nsim,0,sd), nrow=nsim)
> x2 <- matrix(rnorm(n*nsim,diff,sd), nrow=nsim)
> x <- cbind(x1,x2)
> out <- apply(x, 1, function(xx)
+   t.test(xx ~ rep(1:2, each=n))$p.value)
> mean(out <= 0.05)

```

```
[1] 0.9027
```

### Low Power

```

> simfun <- function(n=100, diff=0, sd=1) {
+   x <- rnorm(n, diff, sd)
+   out <- t.test(x)
+   c(p=out$p.value, mean=out$estimate, ci=out$conf.int)
+ }
>
> out <- replicate(1000, simfun(n=50, diff=0.3))
> mean(out[1,] < 0.05)

```


```
[1] 0.553
```

```

> plot(out[2,], seq_len(ncol(out)),
+       xlim=range(out[3:4,]), type='n',
+       ylab='', xlab='Mean')
> segments(x0=out[3,], x1=out[4,], y0=seq_len(ncol(out)),
+          col='green')
> points(out[2,], seq_len(ncol(out)), pch='|')
> abline(v=0.3, col='red')


```





#### Low Power Conditional

```
> w <- out[1,] < 0.05
> plot(out[2,], seq_len(ncol(out)),
+       xlim=range(out[3:4,]), type='n',
+       ylab='', xlab='Mean')
> segments(x0=out[3,], x1=out[4,], y0=seq_len(ncol(out)),
+          col=ifelse(w, 'green', NA))
> points(out[2,], seq_len(ncol(out)), pch=ifelse(w, '|', ''))
> abline(v=0.3, col='red')
```



### Low Power Conditional

```
> mean( out[2,w] > 0.3 )
```

```
[1] 0.8951175
```

```
> mean( out[2,w] > 0.6 )
```

```
[1] 0.03435805
```

```
> mean( out[3,w] > 0.3 )
```

```
[1] 0.05605787
```

### Mixture of Normals

```
> simfun <- function(n=100, diff=0, sd=2, p=0.1) {
+   x <- rnorm(n, diff,
+             ifelse(rbinom(n,1,p), sd, 1))
+   c(t.test(x)$p.value, wilcox.test(x)$p.value)
+ }
> out1 <- replicate(10000, simfun(n=30, diff=0))
> rowMeans(out1 <= 0.05)
```

```
[1] 0.0477 0.0500
```

```
> out2 <- replicate(10000, simfun(n=30, diff=0.5))  
> rowMeans(out2 <= 0.05)
```

```
[1] 0.6549 0.6732
```

### Fishers Exact Test

```
> simfun <- function(n=50, p=c(0.5, 0.5, 0.5)) {  
+   x1 <- sample(1:3, n, replace=TRUE)  
+   x2 <- rbinom(n, 1, p[x1])  
+   fisher.test( table(x1,x2) )$p.value  
+ }  
> out1 <- replicate(10000,  
+                   simfun(n=50, p=c(0.5, 0.5, 0.5)))  
> out2 <- replicate(10000,  
+                   simfun(n=50, p=c(0.5, 0.3, 0.7)))  
> c(mean(out1<=0.05), mean(out2<=0.05))
```

```
[1] 0.0473 0.5188
```

### Beta Binomial CI

```
> simfun <- function(n=100, true.p=0.02, test.p=0.04,  
+                   alpha=0.05) {  
+   x <- rbinom(1, n, true.p)  
+   qbeta(1-alpha/2, x+1, n-x+1) < test.p  
+ }  
> out1 <- replicate(10000, simfun(n=1000, true.p=0.04))  
> mean(out1)
```

```
[1] 0.0171
```

```
> out2 <- replicate(10000, simfun(n=1000))  
> mean(out2)
```

```
[1] 0.9469
```

### Regression

```

> library(MASS)
> simfun <- function(n=100, beta=rep(0,9),
+                   mu.x=rep(0,8), var.x=diag(8),
+                   sig=1) {
+   x <- mvrnorm(n, mu.x, var.x)
+   y <- cbind(1,x) %*% beta + rnorm(n,0,sig)
+   mydat <- cbind( as.data.frame(x), y)
+   names(mydat) <- c(paste0('x',1:8), 'y')
+   fit1 <- lm(y ~ x4 + x5 + x6 + x7 + x8, data=mydat)
+   fit2 <- lm(y ~ ., data=mydat)
+   anova(fit1, fit2)[2,6]
+ }

```

## Logistic Regression

```

> simfun <- function(n=1000, b0=0, b1=0) {
+   x <- rnorm(n)
+   eta <- b0 + b1*x
+   p <- exp(eta)/(1+exp(eta))
+   y <- rbinom(n,1,p)
+   fit <- glm(y~x, family=binomial)
+   coef(summary(fit))[2,4]
+ }

```

## Confidence Interval Width

```

> library(MASS)
> simfun <- function(n=1000, b0=0, b1=0) {
+   x <- rnorm(n)
+   eta <- b0 + b1*x
+   p <- exp(eta)/(1+exp(eta))
+   y <- rbinom(n,1,p)
+   fit <- glm(y~x, family=binomial)
+   ci <- confint(fit,2)
+   c(ci, ci[2]-ci[1])
+ }
>
> out <- replicate(100, simfun(b1=0.8))
> apply(out, 1, range)


```

```

      2.5 %    97.5 %    97.5 %
[1,] 0.4785973 0.7693356 0.2787030
[2,] 0.8181014 1.1545843 0.3370907

```

```
> plot(out[1,], 1:100, type='n', xlim=range(out[1:2,]))
> segments(out[1,],1:100, out[2,])
```



### Cox Model Regression

```
> library(survival)
> simfun <- function(n=100, beta = 0, cens.scale=1) {
+   x <- rnorm(n, 10, 2)
+   y <- rexp(n, 1/(1+beta*x))
+   cens <- rexp(n, 1/cens.scale)
+   time <- pmin(y,cens)
+   status <- ifelse(y<cens, 1, 0)
+   #print(table(status))
+   fit <- coxph(Surv(time,status) ~ x)
+   coef(summary(fit))[1,5]
+ }
```

### Mixed Effects Model

```
> library(lme4)
> simfun <- function(n.student=100, n.school=20,
+                   sig.student=1, sig.school=2,
```

```

+           b0=0, b1=0, b2=0, b12=0) {
+   x1 <- rnorm(n.student*n.school, 0, 1)
+   x2 <- rbinom(n.student*n.school, 1, 0.5)
+   re.school <- rnorm(n.school,0,sig.school)
+   school.id <- rep(1:n.school, each=n.student)
+   y <- b0 + b1*x1 + b2*x2 + b12*x1*x2 +
+     re.school[school.id] +
+     rnorm(n.student*n.school,0,sig.student)
+   fit1 <- lmer( y ~ x1 + (1|school.id))
+   fit2 <- lmer( y ~ x1*x2 + (1|school.id))
+   anova(fit1,fit2)[2,8]
+ }

```

## Multiple Sample Sizes

```

> simfun <- function(n=c(10,20,30,50,75,100), diff=0.4) {
+   x <- rnorm(max(n), diff, 1)
+   sapply(n, function(nn) t.test(x[seq_len(nn)])$p.value)
+ }
>
> out <- replicate(1000, simfun())
> rbind( c(10,20,30,50,75,100),
+   apply(out, 1, function(x) mean(x<=0.05)))

```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	10.000	20.000	30.000	50.000	75.000	100.000
[2,]	0.237	0.403	0.571	0.791	0.926	0.977

## Parallel Processing

```

> library(parallel)
> cl <- makeCluster(4)
> clusterSetRNGStream(cl, 20160405)
> clusterExport(cl, "simfun")
> simfun2 <- function(i,...) simfun(...)
> out <- parSapply(cl, 1:1000, FUN=simfun2, diff=0.6)
> rbind( c(10,20,30,50,75,100),
+   apply(out, 1, function(x) mean(x<=0.05)))

```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	10.000	20.000	30.000	50.000	75	100
[2,]	0.435	0.753	0.908	0.988	1	1

Questions?